

```
In [1]: # Import Libraries
import pandas as pd
import numpy as np
#data visualization
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

Matplotlib is building the font cache; this may take a moment.

```
In [3]: #1 Import the dataset
import os
os.chdir("C:/Users")
```

```
In [4]: #2. Load the dataset into the tool
#add target(age) to dataset [rings+1.5=age]
data=pd.read_csv('abalone.csv')
data['age']=data.Rings+1.5
#remove rings variable
data.drop('Rings',axis=1,inplace=True)
print("Data loaded successfully!")
```

Data loaded successfully!

```
In [5]: df=pd.read_csv('abalone.csv')
```

```
In [6]: df
```

```
Out[6]:
```

|      | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0    | M   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.1500       | 15    |
| 1    | M   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.0700       | 7     |
| 2    | F   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.2100       | 9     |
| 3    | M   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.1550       | 10    |
| 4    | I   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.0550       | 7     |
| ...  | ... | ...    | ...      | ...    | ...          | ...            | ...            | ...          | ...   |
| 4172 | F   | 0.565  | 0.450    | 0.165  | 0.8870       | 0.3700         | 0.2390         | 0.2490       | 11    |
| 4173 | M   | 0.590  | 0.440    | 0.135  | 0.9660       | 0.4390         | 0.2145         | 0.2605       | 10    |
| 4174 | M   | 0.600  | 0.475    | 0.205  | 1.1760       | 0.5255         | 0.2875         | 0.3080       | 9     |
| 4175 | F   | 0.625  | 0.485    | 0.150  | 1.0945       | 0.5310         | 0.2610         | 0.2960       | 10    |
| 4176 | M   | 0.710  | 0.555    | 0.195  | 1.9485       | 0.9455         | 0.3765         | 0.4950       | 12    |

4177 rows × 9 columns

```
In [7]: 5.#check for missing values in the dataset and deal with them
df.isnull().sum()
```

```
Out[7]: Sex          0
Length          0
Diameter        0
Height          0
Whole weight    0
Shucked weight  0
Viscera weight  0
Shell weight    0
Rings          0
dtype: int64
```

```
In [8]: data.describe()
```

```
Out[8]:
```

|              | Length      | Diameter    | Height      | Whole weight | Shucked weight | Viscera weight | Shell weight |
|--------------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|
| <b>count</b> | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000  | 4177.000000    | 4177.000000    | 4177.000000  |
| <b>mean</b>  | 0.523992    | 0.407881    | 0.139516    | 0.828742     | 0.359367       | 0.180594       | 0.238837     |
| <b>std</b>   | 0.120093    | 0.099240    | 0.041827    | 0.490389     | 0.221963       | 0.109614       | 0.139203     |
| <b>min</b>   | 0.075000    | 0.055000    | 0.000000    | 0.002000     | 0.001000       | 0.000500       | 0.001500     |
| <b>25%</b>   | 0.450000    | 0.350000    | 0.115000    | 0.441500     | 0.186000       | 0.093500       | 0.130000     |
| <b>50%</b>   | 0.545000    | 0.425000    | 0.140000    | 0.799500     | 0.336000       | 0.171000       | 0.234000     |
| <b>75%</b>   | 0.615000    | 0.480000    | 0.165000    | 1.153000     | 0.502000       | 0.253000       | 0.329000     |
| <b>max</b>   | 0.815000    | 0.650000    | 1.130000    | 2.825500     | 1.488000       | 0.760000       | 1.005000     |

```
In [9]: data['age'].isnull().sum()
```

```
Out[9]: 0
```

```
In [10]: data['age'].mean()
```

```
Out[10]: 11.433684462532918
```

```
In [11]: data['age'].replace(np.NaN , data['age'].mean()).head(15)
```

```
Out[11]: 0    16.5
1     8.5
2    10.5
3    11.5
4     8.5
5     9.5
6    21.5
7    17.5
8    10.5
9    20.5
10   15.5
11   11.5
12   12.5
13   11.5
14   11.5
Name: age, dtype: float64
```

```
In [12]: data['age'].median()
```

```
Out[12]: 10.5
```

```
In [13]: data['age'].mode()
```

```
Out[13]: 0    10.5
         Name: age, dtype: float64
```

```
In [14]: # 7. Check for categorical columns and perform encoding
         from sklearn.preprocessing import OneHotEncoder
```

```
In [15]: encoder = OneHotEncoder(sparse=False)
         cat_cols = ['sex']
```

```
In [16]: from sklearn.preprocessing import StandardScaler
         # copying original dataframe
         df_ready = df.copy()
```

```
In [17]: from sklearn.preprocessing import StandardScaler
         # copying original dataframe
         df_ready = df.copy()
```

```
In [19]: scaler = StandardScaler()
         num_cols = ['Rings', 'Shell weight', 'Viscera weight', 'Shucked weight', 'Whole we:
```

```
In [20]: df_ready.head()
```

```
Out[20]:
```

|   | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.150        | 15    |
| 1 | M   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.070        | 7     |
| 2 | F   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.210        | 9     |
| 3 | M   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.155        | 10    |
| 4 | I   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.055        | 7     |

```
In [21]: from sklearn.preprocessing import OneHotEncoder
```

```
In [22]: encoder = OneHotEncoder(sparse=False)
         cat_cols = ['Sex']
```

```
In [23]: # Encode Categorical Data
         df_encoded = pd.DataFrame(encoder.fit_transform(df_ready[cat_cols]))
         df_encoded.columns = encoder.get_feature_names(cat_cols)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function get\_feature\_names is deprecated; get\_feature\_names is deprecated in 1.0 and will be removed in 1.2. Please use get\_feature\_names\_out instead.  
warnings.warn(msg, category=FutureWarning)

```
In [24]: # Replace Categorical Data with Encoded Data
         df_ready = df_ready.drop(cat_cols, axis=1)
         df_ready = pd.concat([df_encoded, df_ready], axis=1)
```

```
In [25]: df_ready['Rings'] = df_ready['Rings'].apply(lambda x: 1 if x == 'yes' else 0)
```

```
In [26]: print('Shape of dataframe:', df_ready.shape)
```

```
Shape of dataframe: (4177, 11)
```

In [27]: `df_ready.head()`

Out[27]:

|   | Sex_F | Sex_I | Sex_M | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-------|-------|-------|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | 0.0   | 0.0   | 1.0   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.150        | 0     |
| 1 | 0.0   | 0.0   | 1.0   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.070        | 0     |
| 2 | 1.0   | 0.0   | 0.0   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.210        | 0     |
| 3 | 0.0   | 0.0   | 1.0   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.155        | 0     |
| 4 | 0.0   | 1.0   | 0.0   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.055        | 0     |

In [28]: `# 10.split the data into training and testing`  
`# 12.train the model`  
`# 13.test the model`  
`#Split Dataset for Training and Testing`  
`# Select Features`  
`feature = df_ready.drop('Rings', axis=1)`

In [29]: `# Select Target`  
`target = df_ready['Rings']`

In [30]: `# Set Training and Testing Data`  
`from sklearn.model_selection import train_test_split`  
`X_train, X_test, y_train, y_test = train_test_split(feature , target,`  
`shuffle = True,`  
`test_size=0.2,`  
`random_state=1)`

In [31]: `# Show the Training and Testing Data`  
`print('Shape of training feature:', X_train.shape)`  
`print('Shape of testing feature:', X_test.shape)`  
`print('Shape of training label:', y_train.shape)`  
`print('Shape of testing label:', y_test.shape)`

Shape of training feature: (3341, 10)  
 Shape of testing feature: (836, 10)  
 Shape of training label: (3341,)  
 Shape of testing label: (836,)

In [32]: `X_train`

Out[32]:

|             | Sex_F | Sex_I | Sex_M | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight |
|-------------|-------|-------|-------|--------|----------|--------|--------------|----------------|----------------|--------------|
| <b>666</b>  | 0.0   | 0.0   | 1.0   | 0.455  | 0.350    | 0.120  | 0.4835       | 0.1815         | 0.1440         | 0.1600       |
| <b>2813</b> | 0.0   | 1.0   | 0.0   | 0.255  | 0.195    | 0.055  | 0.0725       | 0.0285         | 0.0170         | 0.0210       |
| <b>1862</b> | 0.0   | 1.0   | 0.0   | 0.520  | 0.410    | 0.110  | 0.5185       | 0.2165         | 0.0915         | 0.1840       |
| <b>3684</b> | 0.0   | 1.0   | 0.0   | 0.620  | 0.470    | 0.155  | 0.9660       | 0.4470         | 0.1710         | 0.2840       |
| <b>551</b>  | 0.0   | 1.0   | 0.0   | 0.615  | 0.490    | 0.155  | 0.9885       | 0.4145         | 0.1950         | 0.3450       |
| ...         | ...   | ...   | ...   | ...    | ...      | ...    | ...          | ...            | ...            | ...          |
| <b>2895</b> | 0.0   | 1.0   | 0.0   | 0.540  | 0.415    | 0.110  | 0.6190       | 0.2755         | 0.1500         | 0.1765       |
| <b>2763</b> | 0.0   | 1.0   | 0.0   | 0.550  | 0.425    | 0.135  | 0.6560       | 0.2570         | 0.1700         | 0.2030       |
| <b>905</b>  | 0.0   | 1.0   | 0.0   | 0.320  | 0.240    | 0.090  | 0.1575       | 0.0700         | 0.0265         | 0.0425       |
| <b>3980</b> | 1.0   | 0.0   | 0.0   | 0.525  | 0.410    | 0.115  | 0.7745       | 0.4160         | 0.1630         | 0.1800       |
| <b>235</b>  | 0.0   | 1.0   | 0.0   | 0.295  | 0.225    | 0.080  | 0.1240       | 0.0485         | 0.0320         | 0.0400       |

3341 rows × 10 columns

In [33]: y\_train

Out[33]:

```

666      0
2813      0
1862      0
3684      0
551       0
..
2895      0
2763      0
905       0
3980      0
235       0
Name: Rings, Length: 3341, dtype: int64

```

In [34]: X\_train.shape

Out[34]: (3341, 10)

In [35]: y\_train.shape

Out[35]: (3341,)

In [36]: X\_train = X\_train.values.reshape((-1,1))

In [37]: X\_train

Out[37]:

```

array([[0.    ],
       [0.    ],
       [1.    ],
       ...,
       [0.0485],
       [0.032 ],
       [0.04  ]])

```

In [38]: y\_train

```
Out[38]: 666      0
          2813    0
          1862    0
          3684    0
          551     0
          ..
          2895    0
          2763    0
          905     0
          3980    0
          235     0
          Name: Rings, Length: 3341, dtype: int64
```

```
In [39]: X_test
```

```
Out[39]:
```

|             | Sex_F | Sex_I | Sex_M | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight |
|-------------|-------|-------|-------|--------|----------|--------|--------------|----------------|----------------|--------------|
| <b>17</b>   | 1.0   | 0.0   | 0.0   | 0.440  | 0.340    | 0.100  | 0.4510       | 0.1880         | 0.0870         | 0.1300       |
| <b>1131</b> | 0.0   | 0.0   | 1.0   | 0.565  | 0.435    | 0.150  | 0.9900       | 0.5795         | 0.1825         | 0.2060       |
| <b>299</b>  | 0.0   | 0.0   | 1.0   | 0.370  | 0.280    | 0.105  | 0.2340       | 0.0905         | 0.0585         | 0.0750       |
| <b>1338</b> | 0.0   | 0.0   | 1.0   | 0.580  | 0.455    | 0.135  | 0.7955       | 0.4050         | 0.1670         | 0.2040       |
| <b>2383</b> | 1.0   | 0.0   | 0.0   | 0.525  | 0.390    | 0.135  | 0.6005       | 0.2265         | 0.1310         | 0.2100       |
| ...         | ...   | ...   | ...   | ...    | ...      | ...    | ...          | ...            | ...            | ...          |
| <b>1787</b> | 0.0   | 1.0   | 0.0   | 0.545  | 0.420    | 0.165  | 0.8935       | 0.4235         | 0.2195         | 0.2280       |
| <b>3075</b> | 1.0   | 0.0   | 0.0   | 0.680  | 0.520    | 0.185  | 1.4940       | 0.6150         | 0.3935         | 0.4060       |
| <b>2766</b> | 1.0   | 0.0   | 0.0   | 0.555  | 0.445    | 0.175  | 1.1465       | 0.5510         | 0.2440         | 0.2785       |
| <b>1410</b> | 1.0   | 0.0   | 0.0   | 0.665  | 0.530    | 0.180  | 1.4910       | 0.6345         | 0.3420         | 0.4350       |
| <b>2529</b> | 1.0   | 0.0   | 0.0   | 0.600  | 0.500    | 0.155  | 1.3320       | 0.6235         | 0.2835         | 0.3500       |


836 rows × 10 columns

```
In [40]: y_test
```

```
Out[40]: 17      0
          1131    0
          299     0
          1338    0
          2383    0
          ..
          1787    0
          3075    0
          2766    0
          1410    0
          2529    0
          Name: Rings, Length: 836, dtype: int64
```


```
In [41]: # 3.perform Visualization
          #Univarient analysis
          sns.displot(df['Sex'])
```

```
Out[41]: <seaborn.axisgrid.FacetGrid at 0x24865593e20>
```




```
In [42]: #Multivariate analysis
corr = df.corr()
sns.heatmap(corr,xticklabels=corr.columns,yticklabels=corr.columns)
```

Out[42]: <AxesSubplot:>



```
In [44]: #Bi-variant analysis
sns.pairplot(df)
```

Out[44]: <seaborn.axisgrid.PairGrid at 0x24869b4b520>



In [45]: `# 4.Descriptive statistics on the dataset`  
`data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Sex         4177 non-null   object
1   Length      4177 non-null   float64
2   Diameter    4177 non-null   float64
3   Height      4177 non-null   float64
4   Whole weight 4177 non-null   float64
5   Shucked weight 4177 non-null   float64
6   Viscera weight 4177 non-null   float64
7   Shell weight 4177 non-null   float64
8   age         4177 non-null   float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
```

In [46]: `data.describe()`



Out[46]:

|       | Length      | Diameter    | Height      | Whole weight | Shucked weight | Viscera weight | Shell weight |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000  | 4177.000000    | 4177.000000    | 4177.000000  |
| mean  | 0.523992    | 0.407881    | 0.139516    | 0.828742     | 0.359367       | 0.180594       | 0.238837     |
| std   | 0.120093    | 0.099240    | 0.041827    | 0.490389     | 0.221963       | 0.109614       | 0.139203     |
| min   | 0.075000    | 0.055000    | 0.000000    | 0.002000     | 0.001000       | 0.000500       | 0.001500     |
| 25%   | 0.450000    | 0.350000    | 0.115000    | 0.441500     | 0.186000       | 0.093500       | 0.130000     |
| 50%   | 0.545000    | 0.425000    | 0.140000    | 0.799500     | 0.336000       | 0.171000       | 0.234000     |
| 75%   | 0.615000    | 0.480000    | 0.165000    | 1.153000     | 0.502000       | 0.253000       | 0.329000     |
| max   | 0.815000    | 0.650000    | 1.130000    | 2.825500     | 1.488000       | 0.760000       | 1.005000     |

```
In [47]: # 6.outlier handling
df = pd.get_dummies(df)
dummy_df = df
```

```
In [50]: import numpy as np
from collections import Counter

def detect_outliers(df, n, features):
    """
    Takes a dataframe df of features and returns a list of the indices
    corresponding to the observations containing more than n outliers according
    to the Tukey method.
    """
    outlier_indices = []

    # iterate over features(columns)
    for col in features:
        # 1st quartile (25%)
        Q1 = np.percentile(df[col], 25)
        # 3rd quartile (75%)
        Q3 = np.percentile(df[col], 75)
        # Interquartile range (IQR)
        IQR = Q3 - Q1

        # outlier step
        outlier_step = 1.5 * IQR

        # Determine a list of indices of outliers for feature col
        outlier_list_col = df[(df[col] < Q1 - outlier_step) | (df[col] > Q3 + outlier_step)]

        # append the found outlier indices for col to the list of outlier indices
        outlier_indices.extend(outlier_list_col)

    # select observations containing more than 2 outliers
    outlier_indices = Counter(outlier_indices)
    multiple_outliers = list(k for k, v in outlier_indices.items() if v > n)

    return multiple_outliers
```

```
In [51]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Length                4177 non-null   float64
1   Diameter              4177 non-null   float64
2   Height                4177 non-null   float64
3   Whole weight          4177 non-null   float64
4   Shucked weight        4177 non-null   float64
5   Viscera weight        4177 non-null   float64
6   Shell weight          4177 non-null   float64
7   Rings                 4177 non-null   int64
8   Sex_F                 4177 non-null   uint8
9   Sex_I                 4177 non-null   uint8
10  Sex_M                 4177 non-null   uint8
dtypes: float64(7), int64(1), uint8(3)
memory usage: 273.4 KB
```

```
In [52]: # 8.split the data into dependent and independent variables
# 9.scale independent variable
# x-independent variable & y-dependent variable
x=df.iloc[:,1]
```

```
In [53]: x
```

```
Out[53]:
```

|      | Length |
|------|--------|
| 0    | 0.455  |
| 1    | 0.350  |
| 2    | 0.530  |
| 3    | 0.440  |
| 4    | 0.330  |
| ...  | ...    |
| 4172 | 0.565  |
| 4173 | 0.590  |
| 4174 | 0.600  |
| 4175 | 0.625  |
| 4176 | 0.710  |

4177 rows × 1 columns

```
In [54]: df=pd.read_csv('abalone.csv')
```

```
In [55]: df
y=df.iloc[:,1:]
```

```
In [56]: y
```

Out[56]:

|      | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0    | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.1500       | 15    |
| 1    | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.0700       | 7     |
| 2    | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.2100       | 9     |
| 3    | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.1550       | 10    |
| 4    | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.0550       | 7     |
| ...  | ...    | ...      | ...    | ...          | ...            | ...            | ...          | ...   |
| 4172 | 0.565  | 0.450    | 0.165  | 0.8870       | 0.3700         | 0.2390         | 0.2490       | 11    |
| 4173 | 0.590  | 0.440    | 0.135  | 0.9660       | 0.4390         | 0.2145         | 0.2605       | 10    |
| 4174 | 0.600  | 0.475    | 0.205  | 1.1760       | 0.5255         | 0.2875         | 0.3080       | 9     |
| 4175 | 0.625  | 0.485    | 0.150  | 1.0945       | 0.5310         | 0.2610         | 0.2960       | 10    |
| 4176 | 0.710  | 0.555    | 0.195  | 1.9485       | 0.9455         | 0.3765         | 0.4950       | 12    |

4177 rows × 8 columns

In [57]: `# 11. Build the model`  
`transformed_sex_feature = OneHotEncoder().fit_transform(df['Sex'].values.reshape(-1, 1))`  
`df_sex_encoded = pd.DataFrame(transformed_sex_feature, columns = ["Sex_"+str(int(i)) for i in range(2)])`  
`df = pd.concat([df, df_sex_encoded], axis=1)`

Input In [57]  
`df_sex_encoded = pd.DataFrame(transformed_sex_feature, columns = ["Sex_"+str(int(i)) for i in range(2)])`  
`df = pd.concat([df, df_sex_encoded], axis=1)`  
`df`  
**SyntaxError: invalid syntax**

In [58]: `df.head()`

Out[58]:

|   | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.150        | 15    |
| 1 | M   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.070        | 7     |
| 2 | F   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.210        | 9     |
| 3 | M   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.155        | 10    |
| 4 | I   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.055        | 7     |

In [59]: `# 14. Measure the performance using Metrics`  
`df['Age'] = df['Rings'] + 1.5`  
`df['Age'].head(5)`

Out[59]:  
0 16.5  
1 8.5  
2 10.5  
3 11.5  
4 8.5  
Name: Age, dtype: float64

In [60]: `'''Sex and Age Visualization'''`

```
plt.figure(figsize = (20,7))
sns.swarmplot(x = 'Sex', y = 'Age', data = df, hue = 'Sex')
sns.violinplot(x = 'Sex', y = 'Age', data = df)
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 56.2% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 52.2% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 58.5% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)

Out[60]: <AxesSubplot:xlabel='Sex', ylabel='Age'>



In [ ]: